

ZAHID IMRAN

AI Engineer — Agentic Systems, RAG, MCP

Manchester, UK · zaahidimraan@gmail.com · zahid-imran.pages.dev · linkedin.com/in/zahid-imran · github.com/zaahidimraan

I build production AI systems that have to be right: document-AI pipelines with deterministic verification, agentic workflows with human-in-the-loop control, and LLM evaluation that keeps models honest. Currently a Software Engineer at POWWR in Manchester, working on document AI, agentic tooling and the cost engineering that makes both viable at scale.

EXPERIENCE

Software Engineer — POWWR

Mar 2026 – Present

- Build production document-AI pipelines that turn unstructured business documents into validated, structured records.
- Design deterministic verification around model output — grounding to source text, plausibility bounds, self-consistency — so extractions can be trusted without additional model calls.
- Work on internal automation and data-reconciliation tooling, and on evaluation practice for AI features heading into production.
- Research cost-efficient inference: token accounting, prompt caching, context compaction, and local/open-weight model cascades.
- Lead AI enablement for engineering colleagues — agentic workflows, model selection, token economics and MCP connectors.

Generative AI Engineer — Sparkix Technologies

Sep 2024 – Dec 2024

- Built GenAI applications with LLMs and RAG; shipped Flask and FastAPI services to Azure and AWS.
- Worked with OpenAI APIs across chatbots, vision-based quality reporting and structured JSON outputs with LangChain.

Generative AI Engineer — Horizon Tech Services

May 2024 – Nov 2024

- Speech enhancement with GAN and Transformer models in PyTorch; selected, trained and fine-tuned sep-former and GAN models.

Deep Learning Intern — GrayHat · Final Year Project

Sep 2023 – May 2024

- Automatic video dubbing backend (Flask) reaching 85% accuracy; fine-tuned OpenAI Whisper on Urdu, cutting word error rate by 4%.

SELECTED PROJECTS

Omni-Channel AI Executive Assistant

A 24/7 personal assistant that answers WhatsApp, Instagram and Messenger — and knows when to wake a human.

- Centralized n8n orchestration layer unifying WhatsApp Business, Instagram Graph API and Facebook Messenger.
- Router agent classifies incoming messages with an LLM; routine inquiries get context-aware replies, high-priority items go to human-in-the-loop review via a private Telegram bot.
- An Observer agent logs interactions to Notion/PostgreSQL and compiles a nightly Daily Activity Briefing; n8n Queue Mode keeps responses sub-second under load.

n8n · Meta APIs · LLM routing · Telegram · PostgreSQL · Notion

Autonomous Recruitment Agent (LLM-as-a-Judge)

Resume screening that agrees with human recruiters 85% of the time — with 40% fewer false positives.

- Judge agent grades candidates against a strict rubric with per-criterion confidence scores, instead of generic summaries.
- Validated on a golden dataset of 100 human-reviewed resumes using Ragas and Arize Phoenix.
- Data ingestion standardized with Model Context Protocol (MCP) servers for Google Drive and Airtable — no brittle API glue.
- The judging core is open source and runnable: rubric grading, position-bias measurement, self-consistency and human-agreement scoring, with 37 tests that pass offline.

LLM-as-a-Judge · Ragas · Arize Phoenix · MCP · Airtable

Privacy-Preserving PII Redaction Pipeline

PII detection and redaction at near-zero inference cost — no document ever leaves the machine.

- Local-first NLP pipeline detecting and redacting personally identifiable information in sensitive documents.
- Replaced cloud APIs with a fine-tuned Microsoft Phi-3 small language model running fully locally.
- 100% data sovereignty by design — built for strict GDPR/privacy environments.

Phi-3 (SLM) · Fine-tuning · Edge AI · GDPR

FURTHER PROJECTS

Agentic Travel Intelligence Dashboard — An agent that queries, filters and ranks hotels on its own — then feeds a live BI dashboard.

Function calling · Google Maps API · ETL · PowerBI

Speech Enhancement with Transformers — Audio restoration at 14 dB SI-SNR and 2.25 PESQ, served in real time.

PyTorch · Transformers · SpeechBrain · Streamlit

Localized Text-to-Speech (XTTS-v2) — A TTS voice that finally sounds right in Asian English accents.

XTTS-v2 · Transfer learning · Audio pipelines

DubLingo — Urdu! Arabic Drama Dubbing — Automating the dubbing of Urdu drama into Arabic, so the work reaches an audience it never had.

LLMs · Whisper · Python · Speech pipelines

Emotion Detection from Voice — Reading emotion from speech at 72% accuracy across four model architectures.

CNN · LSTM · Deep Learning · PyTorch

FOG Prediction Pipeline (MLOps) — A full training-to-deployment loop that runs itself on every commit.

MLflow · DVC · Docker · GitHub Actions

SKILLS

Agentic Orchestration: n8n, LangGraph, Model Context Protocol (MCP), Claude Code, Multi-Agent Systems, ReAct / Plan-and-Execute, Function Calling

GenAI & LLMs: Agentic RAG, LLM-as-a-Judge, Fine-Tuning (LoRA/QLoRA), Structured Outputs, Azure OpenAI, Local LLMs (Ollama), Prompt Compression, Voice Agents

Evaluation & Ops: Ragas, Arize Phoenix, Latency (TTFT) Optimization, LLM Cost Optimization, Prompt Caching, MLflow, DVC

Data & Vector Engineering: Qdrant, Pinecone, PostgreSQL (pgvector), Redis (Semantic Caching), Airtable

Core Development: Python, TypeScript, Next.js, FastAPI, Flask, .NET (C#), Docker, AWS Lambda, PyTorch

EDUCATION

MSc Data Science (Distinction)

Jan 2025 – Mar 2026

University of Salford, Manchester

Advanced Databases · Big Data Tools & Techniques · Machine Learning and Data Mining · Applied Statistics and Data Visualization

BSc Computer Science (GPA 3.16/4)

Aug 2020 – Jun 2024

FAST NUCES, Pakistan

Artificial Intelligence · Statistical Modeling · Deep Learning · MLOps · Generative AI · Cloud Computing · Distributed Data Engineering

CERTIFICATES (16)

Claude Code: A Highly Agentic Coding Assistant — DeepLearning.AI, Feb 2026

Evaluating AI Agents — DeepLearning.AI, Jan 2026

MCP: Build Rich-Context AI Apps with Anthropic — DeepLearning.AI, Feb 2026

AI Agentic Design Patterns with AutoGen — DeepLearning.AI, Jan 2026

Pretraining LLM — DeepLearning.AI, Jan 2026

Prompt Compression and Query Optimization — DeepLearning.AI, Jan 2026

Agentic AI — DeepLearning.AI, Dec 2025

Building Agentic RAG with LlamaIndex — DeepLearning.AI, Dec 2025

Building AI Voice Agents for Production — DeepLearning.AI, Jan 2025

Building Live Voice Agents with Google's ADK — DeepLearning.AI, Dec 2024

Vector Databases: From Embeddings to Applications — DeepLearning.AI, Aug 2024

LangChain for LLM Application Development — DeepLearning.AI, Jul 2024

Preprocessing Unstructured Data for LLM Applications — DeepLearning.AI, Jun 2024

Open Source Models with Hugging Face — DeepLearning.AI, May 2024

AWS Academy: Microservices & CI/CD Pipeline Builder — AWS, Apr 2024

Generative AI with Large Language Models — Coursera, Nov 2023