

ZAHID IMRAN

AI Engineer — Agentic Systems, RAG, MCP

Manchester, UK · zaahidimraan@gmail.com · zahid-imran.pages.dev · linkedin.com/in/zahid-imran · github.com/zaahidimraan

I build production AI systems that have to be right: document-AI pipelines with deterministic verification, agentic workflows with human-in-the-loop control, and LLM evaluation that keeps models honest. Currently an AI Engineer at POWWR in Manchester, working on document AI, agentic tooling and the cost engineering that makes both viable at scale.

EXPERIENCE

AI Engineer — POWWR

Mar 2026 – Present

- Build production document-AI pipelines that turn unstructured business documents into validated, structured records.
- Design deterministic verification around model output — grounding to source text, plausibility bounds, self-consistency — so extractions can be trusted without additional model calls.
- Work on internal automation and data-reconciliation tooling, and on evaluation practice for AI features heading into production.
- Research cost-efficient inference: token accounting, prompt caching, context compaction, and local/open-weight model cascades.
- Lead AI enablement for engineering colleagues — agentic workflows, model selection, token economics and MCP connectors.

Generative AI Engineer — Sparkix Technologies

Sep 2024 – Dec 2024

- Built GenAI applications with LLMs and RAG; shipped Flask and FastAPI services to Azure and AWS.
- Worked with OpenAI APIs across chatbots, vision-based quality reporting and structured JSON outputs with LangChain.

Generative AI Engineer — Horizon Tech Services

May 2024 – Nov 2024

- Speech enhancement with GAN and Transformer models in PyTorch; selected, trained and fine-tuned sep-former and GAN models.

Deep Learning Intern — GrayHat · Final Year Project

Sep 2023 – May 2024

- Automatic video dubbing backend (Flask) reaching 85% accuracy; fine-tuned OpenAI Whisper on Urdu, cutting word error rate by 4%.

SELECTED PROJECTS

Omni-Channel AI Executive Assistant

A 24/7 personal assistant that answers WhatsApp, Instagram and Messenger — and knows when to wake a human.

- Centralized n8n orchestration layer unifying WhatsApp Business, Instagram Graph API and Facebook Messenger.
- Router agent classifies incoming messages with an LLM; routine inquiries get context-aware replies, high-priority items go to human-in-the-loop review via a private Telegram bot.
- An Observer agent logs interactions to Notion/PostgreSQL and compiles a nightly Daily Activity Briefing; n8n Queue Mode keeps responses sub-second under load.

n8n · Meta APIs · LLM routing · Telegram · PostgreSQL · Notion

Autonomous Recruitment Agent (LLM-as-a-Judge)

Resume screening that agrees with human recruiters 85% of the time — with 40% fewer false positives.

- Judge agent grades candidates against a strict rubric with per-criterion confidence scores, instead of generic summaries.
- Validated on a golden dataset of 100 human-reviewed resumes using Ragas and Arize Phoenix.
- Data ingestion standardized with Model Context Protocol (MCP) servers for Google Drive and Airtable — no brittle API glue.
- The judging core is open source and runnable: rubric grading, position-bias measurement, self-consistency and human-agreement scoring, with 37 tests that pass offline.

LLM-as-a-Judge · Ragas · Arize Phoenix · MCP · Airtable

Privacy-Preserving PII Redaction Pipeline

PII detection and redaction at near-zero inference cost — no document ever leaves the machine.

- Local-first NLP pipeline detecting and redacting personally identifiable information in sensitive documents.
- Replaced cloud APIs with a fine-tuned Microsoft Phi-3 small language model running fully locally.
- 100% data sovereignty by design — built for strict GDPR/privacy environments.

Phi-3 (SLM) · Fine-tuning · Edge AI · GDPR

FURTHER PROJECTS

Agentic Travel Intelligence Dashboard — An agent that queries, filters and ranks hotels on its own — then feeds a live BI dashboard.

Function calling · Google Maps API · ETL · PowerBI

Speech Enhancement with Transformers — Audio restoration at 14 dB SI-SNR and 2.25 PESQ, served in real time.

PyTorch · Transformers · SpeechBrain · Streamlit

Localized Text-to-Speech (XTTS-v2) — A TTS voice that finally sounds right in Asian English accents.

XTTS-v2 · Transfer learning · Audio pipelines

DubLingo — Urdu! Arabic Drama Dubbing— Automating the dubbing of Urdu drama into Arabic, so the work reaches an audience it never had.

LLMs · Whisper · Python · Speech pipelines

Emotion Detection from Voice — Reading emotion from speech at 72% accuracy across four model architectures.

CNN · LSTM · Deep Learning · PyTorch

FOG Prediction Pipeline (MLOps) — A full training-to-deployment loop that runs itself on every commit.

MLflow · DVC · Docker · GitHub Actions

SKILLS

Agentic Orchestration: n8n, LangGraph, Model Context Protocol (MCP), Claude Code, Multi-Agent Systems, ReAct / Plan-and-Execute, Function Calling

GenAI & LLMs: Agentic RAG, LLM-as-a-Judge, Fine-Tuning (LoRA/QLoRA), Structured Outputs, Azure OpenAI, Local LLMs (Ollama), Prompt Compression, Voice Agents

Evaluation & Ops: Ragas, Arize Phoenix, Latency (TTFT) Optimization, LLM Cost Optimization, Prompt Caching, MLflow, DVC

Data & Vector Engineering: Qdrant, Pinecone, PostgreSQL (pgvector), Redis (Semantic Caching), Airtable

Core Development: Python, TypeScript, Next.js, FastAPI, Flask, .NET (C#), Docker, AWS Lambda, PyTorch

EDUCATION

MSc Data Science (Distinction)

Jan 2025 – Mar 2026

University of Salford, Manchester

Advanced Databases · Big Data Tools & Techniques · Machine Learning and Data Mining · Applied Statistics and Data Visualization

BSc Computer Science (GPA 3.16/4)

Aug 2020 – Jun 2024

FAST NUCES, Pakistan

Artificial Intelligence · Statistical Modeling · Deep Learning · MLOps · Generative AI · Cloud Computing · Distributed Data Engineering

CERTIFICATES (16)

Claude Code: A Highly Agentic Coding Assistant — DeepLearning.AI, Feb 2026

Evaluating AI Agents — DeepLearning.AI, Jan 2026

MCP: Build Rich-Context AI Apps with Anthropic — DeepLearning.AI, Feb 2026

AI Agentic Design Patterns with AutoGen — DeepLearning.AI, Jan 2026

Pretraining LLM — DeepLearning.AI, Jan 2026

Prompt Compression and Query Optimization — DeepLearning.AI, Jan 2026

Agentic AI — DeepLearning.AI, Dec 2025

Building Agentic RAG with LlamaIndex — DeepLearning.AI, Dec 2025

Building AI Voice Agents for Production — DeepLearning.AI, Jan 2025

Building Live Voice Agents with Google's ADK — DeepLearning.AI, Dec 2024

Vector Databases: From Embeddings to Applications — DeepLearning.AI, Aug 2024

LangChain for LLM Application Development — DeepLearning.AI, Jul 2024

Preprocessing Unstructured Data for LLM Applications — DeepLearning.AI, Jun 2024

Open Source Models with Hugging Face — DeepLearning.AI, May 2024

AWS Academy: Microservices & CI/CD Pipeline Builder — AWS, Apr 2024

Generative AI with Large Language Models — Coursera, Nov 2023